


The 1chipML library



Manual written and maintained by

Jean Michel Sellier*

in collaboration with

**Qinan Qi, Hardik Dalal,
Yimin Nie, Salman Memon.**

*Global Artificial Intelligence Accelerator,
Ericsson, Montréal,
Québec, Canada

manual version 20220510

Contents

1	What is 1chipML?	5
1.1	Introduction	5
1.2	1chipML in a nutshell	6
1.3	Development paradigms	7
1.4	A first example	7
2	Numerical crunching	9
2.1	Random number generators	9
2.1.1	The linear congruential generator	9
2.1.2	The Mersenne twister	10
2.2	Systems of linear equations	10
2.2.1	The Gaussian elimination method	10
2.2.2	The LU decomposition method	11
2.3	Interpolation and extrapolation	11
2.3.1	Polynomial-based approach	12
2.3.2	Spline-based approach	12
2.4	Optimization problems	12
2.4.1	The gradient descent method	12
2.4.2	The genetic approach	12
2.5	Numerical derivation and integration	12
2.5.1	First and second order finite differences approach	12
2.5.2	The Monte Carlo approach	12
2.6	Eigenproblems	14
2.6.1	The Jacobi method	14
2.6.2	The Lanczos method	14
2.7	Statistical approaches	15
2.7.1	Analysis of variance	15
2.7.2	Correlation	15
2.7.3	Cluster analysis	15
2.7.4	Regression	15
2.8	The Fast Fourier Transform	15

3	Machine learning	17
3.1	Neural networks	17
3.1.1	Migration of pre-trained neural networks	17
3.2	Reinforcement Learning	17
3.2.1	The multi-arm bandit problem	17
3.2.2	Monte Carlo methods	17
4	Hardware applications	19
5	How to contribute	21
6	License	23

Chapter 1

What is 1chipML?

1.1 Introduction

The Internet of Things (IoT) and Edge Computing (EC) are, nowadays, topics of high interest since it is becoming clear that important advantages could emerge from this new sort of technologies. What kind of hardware will be in use in this context is still work in progress since many different directions are currently being explored.

In more details, one can consider the IoT as an ensemble of computing objects with sensors, embedded software, etc., which can connect and exchange data with other devices over the Internet (or, equivalently, other communication networks). A continuously growing number of IoT devices are already being created for connected vehicles, home automation, wearable technology, and all sort of appliances with remote monitoring capabilities. For instance, in industry, we are witnessing the creation of a growing number of IoT devices to acquire and analyze data from connected equipment, operational technology, locations, and people. Combined with monitoring devices, this new technological paradigm is helping to regulate and monitor industrial systems. The same approach can be applied for automated record updates of asset placement in industrial storage units. Another important example of IoT is represented by the Internet of Medical Things (IoMT) which can be considered as an application of the IoT for medical and health related purposes. This novel technology can help in the creation of a digitized healthcare system, connecting available medical resources and healthcare services.

Edge computing is a growing and relevant approach as well, but very different than the IoT. In practice, the aim of edge computing is to move the computation away from data centers towards the edge of the network, exploiting smart objects, mobile phones, or network gateways to perform tasks and provide services on behalf of the cloud. Therefore, it also can be considered as a distributed computing paradigm that brings computation and data storage closer to the sources of data. In other words, one can consider the EC as any

type of computer program which can deliver low latency nearer to the requests. Thus, this new computing approach is expected to speed up response times and save bandwidth since it is possible to provide content caching, service delivery, persistent data storage, and IoT management resulting in better response times and transfer rates.

Obviously, the sort of hardware needed to implement these two approaches cannot solely rely on classical computing devices such as servers, Beowulf clusters, etc. since they require those devices to be fast, low power demanding and in small packages. One technology approach that seems to be currently emerging is based on the use of microcontrollers (MCUs). While this seems to be a promising direction, it obviously comes with new challenges which needs to be addressed. Among these open problems, one that seems to be relevant nowadays is the lack of a available software infrastructures for the fast development of such technologies. 1chipML is a library which purpose is to bridge this gap.

1.2 1chipML in a nutshell

In order to make a technology take off, it is important to have a fast and reliable way to develop (and evolve) it. Without such infrastructure and tools, it is hard to imagine how to develop something that can have a real impact on society. In fact, the availability of such infrastructure can become the very reason for the success (or failure) of an invention, no matter what the field is. Obviously, this is valid for the fields of IoT and EC too. At the time of writing this manual, one quickly realizes (unfortunately) that such development tools are practically unavailable, especially if the goal is to run relatively sophisticated algorithms on MCUs such as number crunching and machine learning (ML) methods. Of course, there are some available libraries but they are either proprietary (or incomplete), which is known to slow down the development of technology (a great example is provided by the GNU/Linux operating system which has, eventually, allowed the development of the Android operating system). The goal of 1chipML is to bridge this gap so that IoT and EC can become quickly mainstream.

The target of the 1chipML team of volunteers is to develop and maintain a library which can provide relevant computational capabilities on MCUs related to number crunching and ML. For example, one could need to run the Gauss elimination method to solve a relatively small system of linear equations. The only option for now is to develop everything from scratch which can be time consuming and error prone. With the 1chipML library, this problem is solved. One simply needs to download it and use it right away. A practical example is reported in a section below.

At the time of writing this manual, there are relevant main families of MCUs available on the market such as the ones produced by ARM and AVR. While, in the past, these MCUs used to be very limited (with very limited amount of memory and computational power), nowadays they have features which allow to run relatively complex algorithms on such devices which used to be practically

impossible previously. This is one of the reasons why 1chipML has been created.

1.3 Development paradigms

How to use 1chipML? There are essentially two ways:

- It is possible to include the whole library into a code so that all methods implemented are available at once.
- It is possible to include only the algorithm that is needed. In this case, only the relevant methods are included.

A first practical example is discussed below.

1.4 A first example

The Gauss elimination method, also known as row reduction method, is an algorithm to solve systems of linear equations. This method consists of a sequence of operations performed on the matrix of coefficients corresponding to the system at hand. To perform the computation, a sequence of elementary row operations which modifies the matrix until the lower left-hand corner of the matrix is filled with zeros, as much as possible.

This method is available in the 1chipML library and it is very simple to use it. The user does not need to implement anything related to number crunching beyond the definition of the linear system itself.

```
#include<stdio.h>
#include<stdlib.h>
#include<math.h>

#include " ../src/1chipml.h"

int main(void){
    /* Variables and pointers declarations */
    int i;
    int N;
    gauss_real A[2][2];
    gauss_real B[2];
    gauss_real *sol; /* pointer towards the solution */

    /* The following parameters define a system of 2 linear equations A*X=B */
    N=2;
    A[0][0]=+1.; A[0][1]=+1.;
    A[1][0]=+1.; A[1][1]=-1.;
    B[0]=0.;
    B[1]=1.;
```

```
/* Apply gauss elimination method to solve the system */  
sol=gauss_elimination(N,A,B);  
  
/* Print solution on screen */  
for (i=0;i<N;i++) printf("sol[%d] = %0.3f\n",i,sol[i]);  
return(0);  
}
```


Chapter 2

Numerical crunching

This chapter introduces and discusses the various numerical methods that are already implemented in the current version of the library. The Reader should bear in mind that this is still work in progress at the moment.

2.1 Random number generators

The generation of random numbers is a process used to create a sequence of numbers which cannot be reasonably predicted. Two main methods are used to generate random numbers. In the first approach, one measures some physical phenomenon that is expected to be random and then compensates for possible biases in the measurement process. For instance, sources of randomness include measuring atmospheric noise, thermal noise, and other external electromagnetic and quantum phenomena. In the second approach, one uses computational algorithms that can produce long sequences of apparently random results, which are in fact completely determined by a shorter initial value, known as a seed value or key. This type of random number generator is often called a pseudorandom number generator and it is the method used in the 1chipML library.

While a pseudorandom number generator based solely on deterministic logic can never be regarded as a true random number source in the purest sense of the word, in practice they are generally sufficient even for demanding security-critical applications. In fact, carefully designed and implemented pseudorandom number generators can be certified for security-critical cryptographic purposes. Many computational methods exist for pseudorandom number generation. In the following we present two of these pseudorandom number generators.

2.1.1 The linear congruential generator

This method is implemented in the file "src/linear_congruential_random_generator.c" of the library.

A linear congruential generator is an algorithm that yields a sequence of

pseudo-randomized numbers computed with a discontinuous piecewise linear equation. This method represents one of the oldest and best-known pseudorandom number generator algorithms. The theory behind them is relatively easy and, consequently, allow a fast and easy implementation.

The main idea behind this generator is to use the following recursive relation:

$$X_{(n+1)} = (aX_n + c) \pmod{m},$$

where $X_{(n+1)}$ and X_n are values in the sequence of pseudorandom numbers, and X_0 is called the seed or start value of the sequence. The constants a , c and m are known as the multiplier, the increment and the modulus respectively with $m > 0$, $0 < a < m$ and $0 \leq c < m$. In our particular implementation of this method, the values for these constants are $a = 1027$, $c = 0$ and $m = 1048576$. The initial seed X_0 is set to 38467 by default. The user can modify it by using the following command:

```
ISEED = some_value;
```

For clarity, an extract of the file "tests/test_linear_congruential_random_generator.c" is reported below:

```
/* Variables and pointers declarations */
int i, n=100;

printf("linear_congruential_random_generator\n");
for(i=0; i<n; i++) printf("%1.3f ", linear_congruential_random_generator());
```

2.1.2 The Mersenne twister

Work in progress. This method is not implemented in the library yet.

2.2 Systems of linear equations

A system of linear equations is simply a collection of one or more linear equations involving the same variables. The theory of linear systems is the basis and a fundamental part of linear algebra, a subject which is used in most parts of modern mathematics. Computational algorithms for finding the solutions are an important part of numerical linear algebra, and play a prominent role in engineering, physics, chemistry, computer science, and economics. In the following we present the methods implemented in the library.

2.2.1 The Gaussian elimination method

This method is implemented in the file "src/gauss_elimination.c" of the library.

The Gaussian elimination method, also known as row reduction, is an algorithm used to solve systems of linear equations. It consists of a sequence of

operations, i.e. row reductions, performed on the corresponding matrix of coefficients. This method can also be used to compute the rank of a matrix, the determinant of a square matrix, and the inverse of an invertible matrix.

To perform row reduction on a matrix, one uses a sequence of elementary row operations to modify the matrix. There are three types of elementary row operations, i.e. 1) swapping two rows, 2) multiplying a row by a nonzero number, and 3) adding a multiple of one row to another row. Using these operations, a matrix can always be transformed into an upper triangular matrix, and in fact one that is in row echelon form. Once all of the leading coefficients are equal to 1, and every column containing a leading coefficient has zeros elsewhere, the matrix is said to be in reduced row echelon form. This final form is unique, i.e. it is independent of the sequence of row operations used, and it is this form that is utilized to find the solutions of the system at hand.

To better understand how to use this method implemented in the library, an extract of the file "tests/test_gauss_elimination.c" is reported below:

```
/* Variables and pointers declarations */
int i;
int N;
gauss_real A[2][2];
gauss_real B[2];
gauss_real *sol; /* pointer towards the solution */

/* The following parameters define a system of 2 linear equations A*X=B */
N=2;
A[0][0]=+1.; A[0][1]=+1.;
A[1][0]=+1.; A[1][1]=-1.;
B[0]=0.;
B[1]=1.;

/* Apply gauss elimination method to solve the system */
sol=gauss_elimination(N,A,B);

/* Print solution on screen */
for (i=0;i<N;i++) printf(" sol[%d] = %0.3f\n",i , sol[i]);
return(0);
```

2.2.2 The LU decomposition method

Work in progress. This method is not implemented in the library yet.

2.3 Interpolation and extrapolation

Work in progress.

2.3.1 Polynomial-based approach

Work in progress.

2.3.2 Spline-based approach

Work in progress.

2.4 Optimization problems

Work in progress.

2.4.1 The gradient descent method

Work in progress.

2.4.2 The genetic approach

Work in progress.

2.5 Numerical derivation and integration

Work in progress.

2.5.1 First and second order finite differences approach

Work in progress.

2.5.2 The Monte Carlo approach

The file "src/mc.integration.c" describes the function of doing integration with Monte Carlo approach. One can examine the expected value of an integral using the Monte Carlo approach. Traditionally, the expected value of a function $g(x)$ can be calculated by first multiplying by its probability density function, $f(x)$, and taking the integral over the desired region:

$$E[g(x)] = \int_a^b g(x)f(x)dx$$

Alternatively, we can use a Monte Carlo approximation for the expected value by repeatedly sampling a uniform distribution between the limits of integration.

$$E[g(x)] = \frac{1}{n} \sum_{i=1}^n f(x_i)$$

where $x_i \in [a, b]$. As noted, x_i is a value that is sampled from a uniform distribution between the limits a and b for each unique $n = 1, 2, 3$, etc. This

approach samples the $f(x)$ function and uses the law of large numbers to find a converged expected value. As an aside, the multiplicative factor of $1/n$ is sometimes given as $1/(n-1)$ because there are truly $n-1$ degrees of freedom with n -samples, but when n is large enough the difference between $1/n$ and $1/(n-1)$ is negligible.

Given the form of the estimator for the expected value, extending to the estimate of the integral is simple. The expected value formula is multiplied by the range of the integration limits, as shown below.

$$F = (b-a) \frac{1}{n} \sum_{i=1}^n f(x_i)$$

where $x_i \in [a, b]$.

The file "src/mc_std.c" describes how to calculate variance using the Monte Carlo approach. The variance of the Monte Carlo integration scheme follows a traditional process of calculating variance around some random variable. By continuing with the notation of taking the integral of a function $g(x)$, and the expected value of the integral being $E[g(x)]$, the relationship for standard deviation can be given as:

$$\sigma_n = V \sqrt{\frac{E[g(x)^2] - E[g(x)]^2}{n-1}}.$$

Here the additional V term represents the total volume of the integration limits. If we were to evaluate a one-dimensional integral between the limits of a to b then the equation could be written as:

$$\sigma_n = (b-a) \sqrt{\frac{E[g(x)^2] - E[g(x)]^2}{n-1}}.$$

Using this format, we can easily calculate standard deviation or variance (standard deviation² = variance) along with the integration estimate.

The file "src/mc_stratified_sampling.c" is concerned with stratified sampling which is an approach to variance reduction by which the integration volume is divided into subdomains that are each evaluated separately. The estimates from each subdomain are then combined with a weight depending on their subdomain integration volume. Stratified sampling has the benefit of always reducing the sample variance, without any prior knowledge of the function's shape. Optimized stratified sampling routines may recursively divide the function into regions with comparable values by grouping specific peaks or values. Just as well, the naive approach of parsing the integration volume into uniformly large works without any prior knowledge of the function's structure. The stratified sampling estimate of an integral over a function $g(x)$ is given as:

$$E[g(x)] = \sum_{j=1}^k \frac{V_j}{n_j} \sum_{i=1}^{n_j} g(x_{ij})$$

where $x_i \in V_j$.

2.6 Eigenproblems

Work in progress.

2.6.1 The Jacobi method

Work in progress.

2.6.2 The Lanczos method

The lanczos method is an iterative algorithm for approximating the eigenvalues and eigenvectors of a symmetrical, sparse matrix.

The algorithm generates a sequence of orthonormal vectors, called the Lanczos basis, that form an approximation of the eigenvectors of the matrix. At each iteration, the algorithm computes values forming a tridiagonal matrix that is similar to the original matrix and has the same eigenvalues.

The basic idea behind the Lanczos algorithm is to generate an orthonormal basis for the subspace spanned by the approximate eigenvectors of the matrix and then use this basis to construct the tridiagonal matrix.

Here is a pseudo-code of the algorithm :

Algorithm 1 Lanczos Algorithm

```

1:  $x_0 \leftarrow$  Random unit vector
2:  $\beta_0 \leftarrow 0$ 
3: for  $i = 1$ ;  $i \leq k$ ;  $i++$  do                                 $\triangleright k$  is the number of iterations
4:    $v_i \leftarrow Ax_{i-1} - \beta_i v_{i-1}$ 
5:    $\alpha_i \leftarrow v_i^T \cdot x_{i-1}$ 
6:    $v_i \leftarrow v_i / \|v_i\|$ 
7:    $x_i \leftarrow v_i$ 
8:    $\beta_i \leftarrow \|v_i\|$ 
9: end for

```

$$T = \begin{bmatrix} \alpha_1 & \beta_1 & \dots & 0 \\ \beta_1 & \alpha_2 & \beta_2 & \vdots \\ \vdots & \beta_2 & \ddots & \beta_k \\ 0 & \dots & \beta_k & \alpha_k \end{bmatrix}$$

$$V = [v_0 \quad v_1 \quad \dots \quad v_n]$$

The Lanczos algorithm is useful for solving large sparse eigenvalue problems since it generates a small, dense tridiagonal matrix that can be solved efficiently, rather than solving the entire matrix. It can also be used to compute a few eigenvalues and eigenvectors of the matrix, rather than all of them.

Due to the rounding error caused by the IEEE-754 standard, the Lanczos algorithm is vulnerable to numerical instability. When applying the algorithm,

Lanczos bases tend to lose their orthogonality. To overcome this issue, the Gram-Schmidt algorithm is applied after every iteration to reorthogonalize the vectors v_i .

The Lanczos algorithm does not directly output the eigenvalues and the eigenvectors associated with the original matrix. Further steps are needed to compute those values. However, the Lanczos algorithm does help finding those values by providing a tridiagonal matrix T which is much easier to compute.

The eigenvalues and eigenvectors of T can be calculated using standard methods such as the Jacobi algorithm.

Once the eigenvectors and eigenvalues of T are known, the eigenvectors of the original matrix can be approximated by linear combinations of the Lanczos basis vectors. For example, given an eigenvector of T , say γ , corresponding eigenvector of the original matrix, say λ is given by $\lambda = V\gamma$

2.7 Statistical approaches

Work in progress.

2.7.1 Analysis of variance

Work in progress.

2.7.2 Correlation

Work in progress.

2.7.3 Cluster analysis

Work in progress.

2.7.4 Regression

Work in progress.

2.8 The Fast Fourier Transform

The Fast Fourier Transform (FFT) is an algorithm that computes the discrete Fourier transform (DFT) of a set of complex values. More specifically, the FFT is often used to convert a set of values from its original domain to a representation in the frequency domain.

The DFT equation can be seen in 2.1:

$$X_k = \sum_{n=0}^{N-1} x_n \exp\left(-\frac{2\pi i}{N}nk\right) \quad (2.1)$$

where N is the size of set of values, k ranges from 0 to $N - 1$ and i is the imaginary part. The equation shown in 2.2 represents the twiddle factor.

$$\exp\left(-\frac{2\pi ik}{N}\right) \quad (2.2)$$

The Cooley–Tukey FFT algorithm is used to find the FFT of a set of values. This algorithm uses a strategy similar to divide-and-conquer, allowing it to have a time complexity of $O(n \log n)$.

The algorithm is implemented in a function that takes in two separate arrays to represent complex values. The first array consists of real numbers, while the second array holds imaginary numbers. This algorithm overrides the arrays passed as arguments and replaces them for the result of the FFT. Furthermore, the radix-2 decimation-in-time (DIT) method was used, meaning the arrays must have a length equal to a power of 2. If this is not the case, an error code will be returned by the function.

The first step of the algorithm is to execute bit-reversal permutation on the incoming arrays, as it uses the iterative approach. Then, the divide-and-conquer strategy starts. For the first set, the algorithm iterates over the elements that require the same twiddle factor. It then iterates over each group and uses the same twiddle factor for each one. After that, the twiddle factor is updated using the trigonometric recurrence formula and the next elements are selected. The same steps are repeated for the second set and so forth. By using this strategy, we end up with the result of the FFT.

Given the context of microcontrollers and their lack of memory, multiple design choices were made, which are summarized here:

- The FFT takes in two arrays of the same size (real and imaginary)
- The FFT only takes in arrays that have a length equal to a power of 2
- The FFT overrides the incoming arrays with the result
- The FFT returns 1 in case of an error, 0 otherwise
- The FFT has a space complexity of $O(1)$ and a time complexity of $O(n \log n)$
- The FFT uses the iterative approach
- The FFT uses bit-reversal for the iterative approach
- The FFT does not calculate the twiddle factors in advance
- The FFT uses the trigonometric recurrence formula to reduce the number of sin and cos calculated

Chapter 3

Machine learning

The various already implemented machine learning methods are described here.

3.1 Neural networks

3.1.1 Migration of pre-trained neural networks

3.2 Reinforcement Learning

3.2.1 The multi-arm bandit problem

ϵ -greedy method, UCB (upper confidence bound) method, gradient bandit method.

3.2.2 Monte Carlo methods

On-policy methods, Off-policy methods.

Chapter 4

Hardware applications

Chapter 5

How to contribute

Chapter 6

License

Apache License
Version 2.0, January 2004
<http://www.apache.org/licenses/>

TERMS AND CONDITIONS FOR USE, REPRODUCTION, AND DISTRIBUTION

1. Definitions.

"License" shall mean the terms and conditions for use, reproduction, and distribution as defined by Sections 1 through 9 of this document.

"Licensor" shall mean the copyright owner or entity authorized by the copyright owner that is granting the License.

"Legal Entity" shall mean the union of the acting entity and all other entities that control, are controlled by, or are under common control with that entity. For the purposes of this definition, "control" means (i) the power, direct or indirect, to cause the direction or management of such entity, whether by contract or otherwise, or (ii) ownership of fifty percent (50%) or more of the outstanding shares, or (iii) beneficial ownership of such entity.

"You" (or "Your") shall mean an individual or Legal Entity exercising permissions granted by this License.

"Source" form shall mean the preferred form for making modifications, including but not limited to software source code, documentation source, and configuration files.

"Object" form shall mean any form resulting from mechanical

transformation or translation of a Source form, including but not limited to compiled object code, generated documentation, and conversions to other media types.

"Work" shall mean the work of authorship, whether in Source or Object form, made available under the License, as indicated by a copyright notice that is included in or attached to the work (an example is provided in the Appendix below).

"Derivative Works" shall mean any work, whether in Source or Object form, that is based on (or derived from) the Work and for which the editorial revisions, annotations, elaborations, or other modifications represent, as a whole, an original work of authorship. For the purposes of this License, Derivative Works shall not include works that remain separable from, or merely link (or bind by name) to the interfaces of, the Work and Derivative Works thereof.

"Contribution" shall mean any work of authorship, including the original version of the Work and any modifications or additions to that Work or Derivative Works thereof, that is intentionally submitted to Licensor for inclusion in the Work by the copyright owner or by an individual or Legal Entity authorized to submit on behalf of the copyright owner. For the purposes of this definition, "submitted" means any form of electronic, verbal, or written communication sent to the Licensor or its representatives, including but not limited to communication on electronic mailing lists, source code control systems, and issue tracking systems that are managed by, or on behalf of, the Licensor for the purpose of discussing and improving the Work, but excluding communication that is conspicuously marked or otherwise designated in writing by the copyright owner as "Not a Contribution."

"Contributor" shall mean Licensor and any individual or Legal Entity on behalf of whom a Contribution has been received by Licensor and subsequently incorporated within the Work.

2. Grant of Copyright License. Subject to the terms and conditions of this License, each Contributor hereby grants to You a perpetual, worldwide, non-exclusive, no-charge, royalty-free, irrevocable copyright license to reproduce, prepare Derivative Works of, publicly display, publicly perform, sublicense, and distribute the Work and such Derivative Works in Source or Object form.
3. Grant of Patent License. Subject to the terms and conditions of this License, each Contributor hereby grants to You a perpetual, worldwide, non-exclusive, no-charge, royalty-free, irrevocable (except as stated in this section) patent license to make, have made,

use, offer to sell, sell, import, and otherwise transfer the Work, where such license applies only to those patent claims licensable by such Contributor that are necessarily infringed by their Contribution(s) alone or by combination of their Contribution(s) with the Work to which such Contribution(s) was submitted. If You institute patent litigation against any entity (including a cross-claim or counterclaim in a lawsuit) alleging that the Work or a Contribution incorporated within the Work constitutes direct or contributory patent infringement, then any patent licenses granted to You under this License for that Work shall terminate as of the date such litigation is filed.

4. Redistribution. You may reproduce and distribute copies of the Work or Derivative Works thereof in any medium, with or without modifications, and in Source or Object form, provided that You meet the following conditions:
 - (a) You must give any other recipients of the Work or Derivative Works a copy of this License; and
 - (b) You must cause any modified files to carry prominent notices stating that You changed the files; and
 - (c) You must retain, in the Source form of any Derivative Works that You distribute, all copyright, patent, trademark, and attribution notices from the Source form of the Work, excluding those notices that do not pertain to any part of the Derivative Works; and
 - (d) If the Work includes a "NOTICE" text file as part of its distribution, then any Derivative Works that You distribute must include a readable copy of the attribution notices contained within such NOTICE file, excluding those notices that do not pertain to any part of the Derivative Works, in at least one of the following places: within a NOTICE text file distributed as part of the Derivative Works; within the Source form or documentation, if provided along with the Derivative Works; or, within a display generated by the Derivative Works, if and wherever such third-party notices normally appear. The contents of the NOTICE file are for informational purposes only and do not modify the License. You may add Your own attribution notices within Derivative Works that You distribute, alongside or as an addendum to the NOTICE text from the Work, provided that such additional attribution notices cannot be construed as modifying the License.

You may add Your own copyright statement to Your modifications and may provide additional or different license terms and conditions for use, reproduction, or distribution of Your modifications, or for any such Derivative Works as a whole, provided Your use, reproduction, and distribution of the Work otherwise complies with the conditions stated in this License.

5. **Submission of Contributions.** Unless You explicitly state otherwise, any Contribution intentionally submitted for inclusion in the Work by You to the Licensor shall be under the terms and conditions of this License, without any additional terms or conditions. Notwithstanding the above, nothing herein shall supersede or modify the terms of any separate license agreement you may have executed with Licensor regarding such Contributions.
6. **Trademarks.** This License does not grant permission to use the trade names, trademarks, service marks, or product names of the Licensor, except as required for reasonable and customary use in describing the origin of the Work and reproducing the content of the NOTICE file.
7. **Disclaimer of Warranty.** Unless required by applicable law or agreed to in writing, Licensor provides the Work (and each Contributor provides its Contributions) on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied, including, without limitation, any warranties or conditions of TITLE, NON-INFRINGEMENT, MERCHANTABILITY, or FITNESS FOR A PARTICULAR PURPOSE. You are solely responsible for determining the appropriateness of using or redistributing the Work and assume any risks associated with Your exercise of permissions under this License.
8. **Limitation of Liability.** In no event and under no legal theory, whether in tort (including negligence), contract, or otherwise, unless required by applicable law (such as deliberate and grossly negligent acts) or agreed to in writing, shall any Contributor be liable to You for damages, including any direct, indirect, special, incidental, or consequential damages of any character arising as a result of this License or out of the use or inability to use the Work (including but not limited to damages for loss of goodwill, work stoppage, computer failure or malfunction, or any and all other commercial damages or losses), even if such Contributor has been advised of the possibility of such damages.
9. **Accepting Warranty or Additional Liability.** While redistributing the Work or Derivative Works thereof, You may choose to offer, and charge a fee for, acceptance of support, warranty, indemnity, or other liability obligations and/or rights consistent with this

License. However, in accepting such obligations, You may act only on Your own behalf and on Your sole responsibility, not on behalf of any other Contributor, and only if You agree to indemnify, defend, and hold each Contributor harmless for any liability incurred by, or claims asserted against, such Contributor by reason of your accepting any such warranty or additional liability.

END OF TERMS AND CONDITIONS

APPENDIX: How to apply the Apache License to your work.

To apply the Apache License to your work, attach the following boilerplate notice, with the fields enclosed by brackets "[]" replaced with your own identifying information. (Don't include the brackets!) The text should be enclosed in the appropriate comment syntax for the file format. We also recommend that a file or class name and description of purpose be included on the same "printed page" as the copyright notice for easier identification within third-party archives.

Copyright [yyyy] [name of copyright owner]

Licensed under the Apache License, Version 2.0 (the "License");
you may not use this file except in compliance with the License.
You may obtain a copy of the License at

<http://www.apache.org/licenses/LICENSE-2.0>

Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License.

Bibliography

- [1] W.S. McCulloch, W. Pitts, A logical calculus of the ideas immanent in nervous activity, The bulletin of mathematical biophysics 5.4, pp. 115-133, (1943).
- [2] F. Rosenblatt, The perceptron: a probabilistic model for information storage and organization in the brain, Psychological review 65.6, p. 386, (1958).
- [3] C.M. Bishop, Neural Networks for Pattern Recognition, Oxford University Press, (1995).
- [4] I. Goodfellow, Y. Bengio, A. Courville, Deep Learning, The MIT Press, (2016).
- [5] H.W. Press, B.P. Flannery, P. Brian, S.A. Teukolsky, W.T. Vetterling, Numerical Recipes: The Art of Scientific Computing (third edition), Cambridge University Press, (2007).